// THE ARES SETUP GUIDE · START HERE

How I'd build an AI chief of staff from scratch

For about 10 months I've run my work through one AI system. It handles my inbox, calendar, content, code and servers, and it remembers everything that matters. This is the version I'd build if I started again today: the same five ideas, without my 10 months of mistakes.

STEP 0 · BEFORE ANYTHING

Give it a home and a voice

An assistant that only exists in a browser tab forgets you when you close it. It needs two things: a machine that's always on (a spare Mac mini or a small cloud server is enough), and a chat app you already live in, like WhatsApp, Telegram or Discord, so you can talk to it from your phone.

> **Do this:** pick the machine and the chat app. That's the whole step.

1 Give it a memory

This is the step that matters most. Keep a folder of plain text files the AI reads at the start of every conversation: who you are, how you like to work, your projects, your people. The key rule: when it learns something important, it writes it down.

> **Do this:** write one page about yourself. Your role, your goals, how you like answers (short? detailed?), and what it should never do.

2 Give it hands

A chatbot answers. A chief of staff acts. Connect your tools one at a time: calendar, email, files, then the web. When it does the same task twice, have it write a small script so the third time is instant. Mine has written about 440 of its own.

> **Do this:** connect your calendar and email, read-only at first.

3 Put it on a schedule

The real shift is when it works without being asked. Mine runs 94 scheduled jobs now, but start with three routines, not ninety:

- a morning brief: today's calendar, what's urgent, what's waiting on you
- an evening wrap-up: what got done, what slipped
- an inbox sweep: sort, flag, draft replies

> **Do this:** set up the morning brief first. You'll know within a week whether it's working.

4 Make it learn from mistakes

It will get things wrong. The difference is whether it gets them wrong twice. Keep a rules file. Every time it messes up, add one line saying what to do next time, and why. After a few months, that file is worth more than any prompt you could buy.

> **Do this:** start the file today, even if it's empty.

5 Keep the final say

Let it draft anything. Don't let it send, post, pay or delete without asking you first. It shows you what it wants to do, and you reply yes or no. That one rule is what makes it safe to give it real access.

> **Do this:** make "draft only, ask before sending" the first line of your rules file.

THE STACK I ACTUALLY USE

The brain	Claude, on a flat subscription rather than pay-per-use, for reasoning, writing and decisions.
Always-on runtime	OpenClaw (open source). It keeps the agent running 24/7 and connects it to WhatsApp and Discord.
Memory	Plain text files, plus GBrain (open source) as a searchable knowledge graph once it grows. Mine holds about 6,700 pages.
Free local work	Ollama and Whisper on the Mac mini: transcribing voice notes and summarising recordings, privately and for free.
Connecting machines	Tailscale, a private network so it can manage several computers safely. I run seven.

YOUR FIRST WEEK

DAY 1–2	Set up the always-on machine and connect one chat app. Send it a message from your phone.
DAY 3	Write your one-page profile and your rules file.
DAY 4	Connect calendar and email, read-only.
DAY 5	Turn on the morning brief.
DAY 6–7	Use it for everything. Every time it's wrong, add a line to the rules file.

DON'T AUTOMATE THESE (YET)

- Sending money or making purchases
- Posting anything publicly
- Emailing or messaging people on your behalf without your OK
- Deleting files, emails or anything else you can't get back

MISTAKES I MADE SO YOU DON'T HAVE TO

- **Speed over accuracy.** On its fastest, cheapest settings it confidently recommended a restaurant based on a 9-month-old listing. Let it think. Make it say how sure it is.
- **Too many alerts.** If everything pings you, you stop reading them. Alerts should only repeat when something gets worse.
- **"Done" isn't proof.** Make it show you the result (the file, the sent email, the live page), not just say it finished.
- **Stale approvals.** If hours pass between your "yes" and the send, make it check the thread again first. You may have already sent it yourself.

You think and decide.
It remembers, executes and follows up on everything else.

Ready to build it? The OpenClaw Bible has the full setup, the receipts and the cost breakdown: github.com/rushindrasinha/openclaw-guide